

Operating Systems - 2

↳ Load Balancers

→ Round Robin CPU Scheduling

→ Solⁿ to starvation

→ Threads → Throughput v/s latency

Theory

→ Intro

→ Multicore v/s Single Core CPU

→ Concurrency v/s Parallelism

→ Code:

* Hello World

* Print 1 to 100

↳ Practical

→ Java

ROUND ROBIN CPU SCHEDULING ALGO

→ FCFS
→ SJF



P1



P2



⇒ Most of the processes will be making no progress

→ We want every program that is running to make some progress

RR

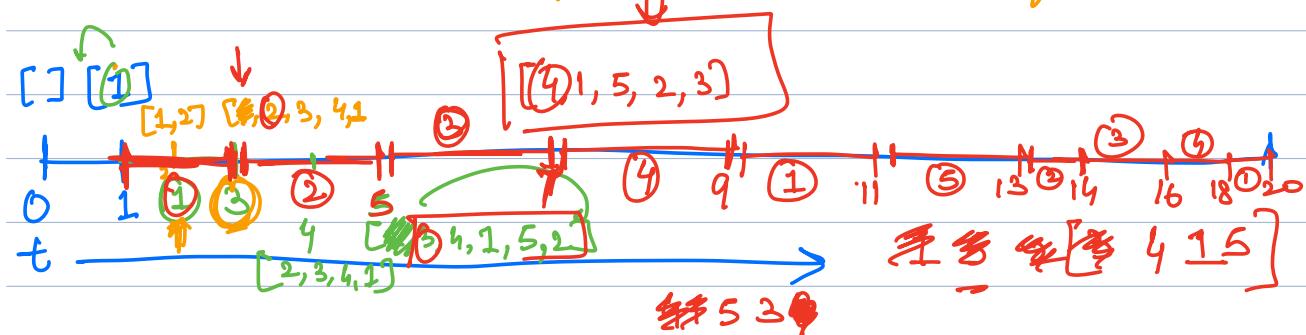
$$q = \text{time quantum} \\ = 2\text{s} \Rightarrow 20\text{sec} \Rightarrow \text{FCFS}$$

$\Rightarrow 0.2\text{ sec}$ \Rightarrow lot of Context Switching

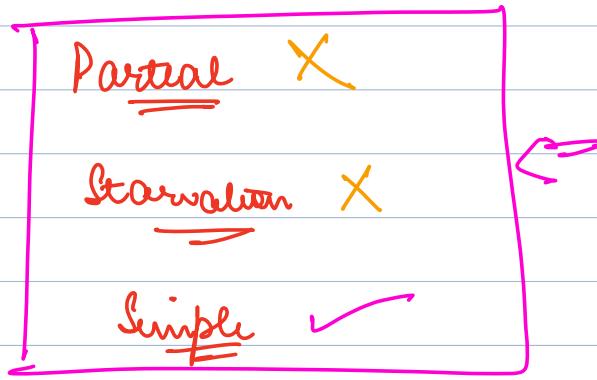
After every q sec, it will schedule some other process to run.

| Process | Arrival Time | Time To Complete |
|---------|--------------|------------------|
| 1 | 1 | 6 4 2 0 |
| 2 | 2 | 8 1 |
| 3 | 3 | 9 7 5 |
| 4 | 3 | 4 2 0 |
| 5 | 5 | 11 9 |

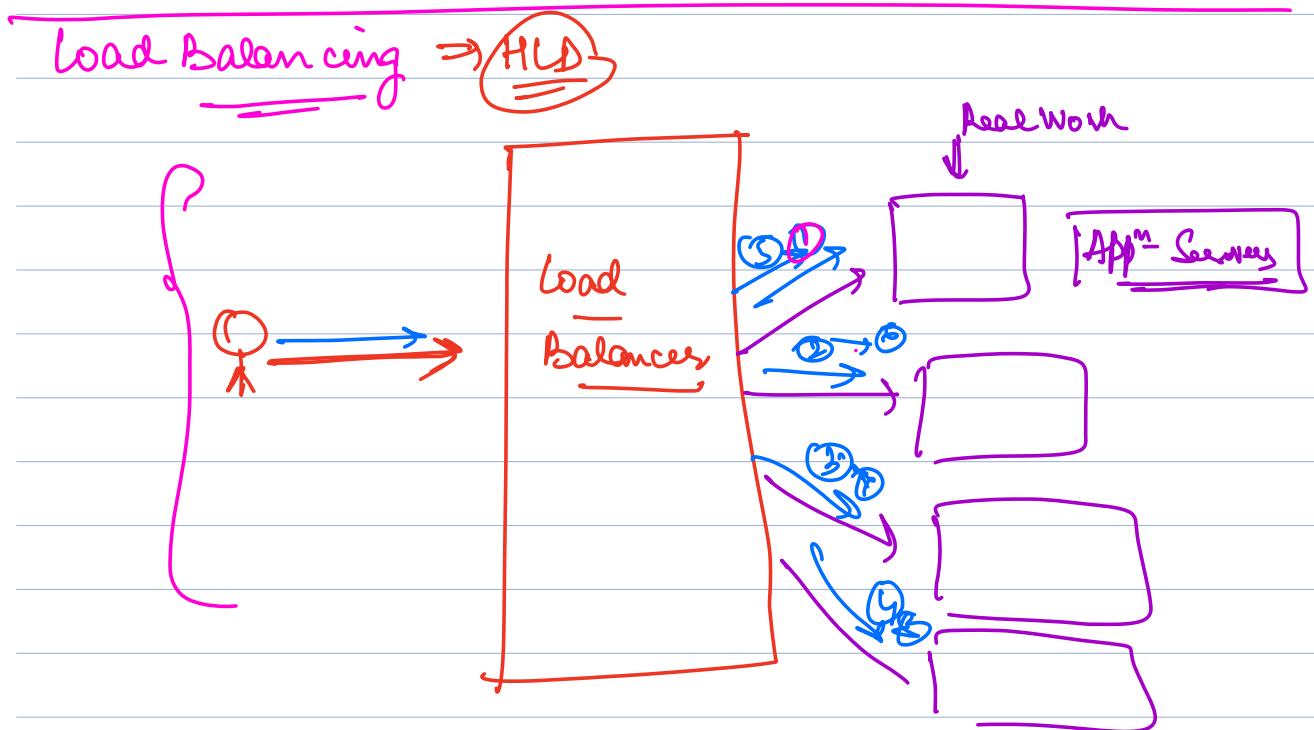
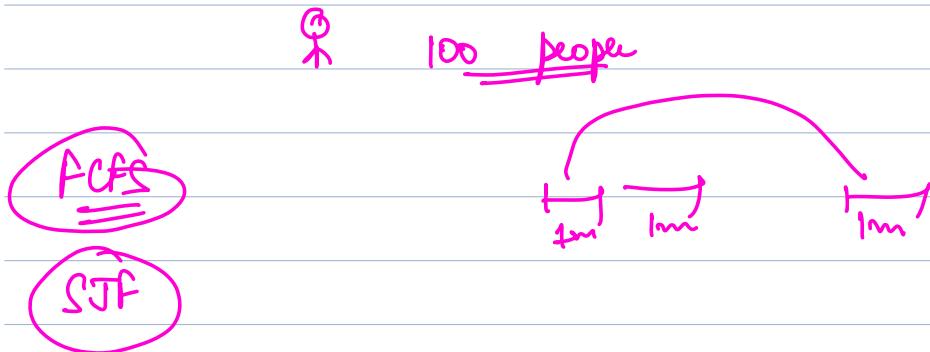
- Maintains a queue []
- Whenever a new process enters, it goes to the back of queue
- Whenever a process finishes ' q ' sec of run, it is put at the back of the queue



~~4, 1, 5, 2, 3~~



Eg Party is happening at home



Balances the load
to individual
servers

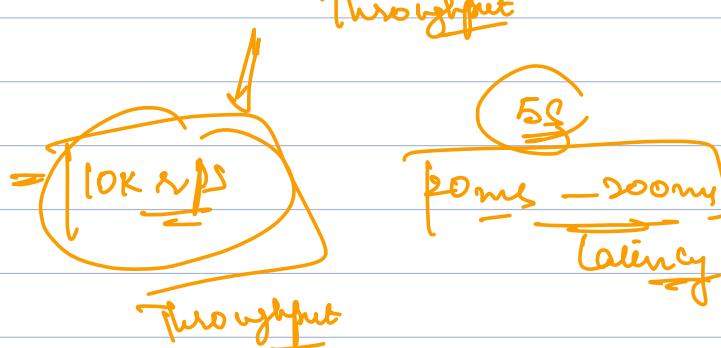
Throughput and latency

~~Throughput~~ \Rightarrow Efficiency
 $\qquad\qquad\qquad$ (Average) $\Rightarrow \left\{ \begin{array}{l} \text{\# processes that are} \\ \text{completed per unit time} \end{array} \right.$

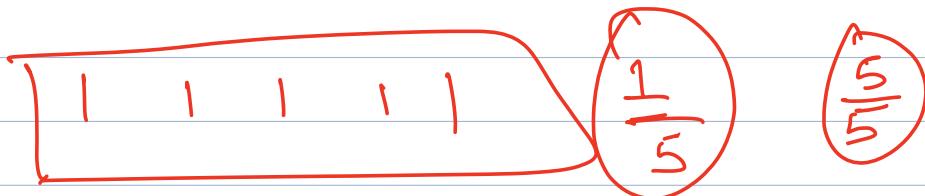
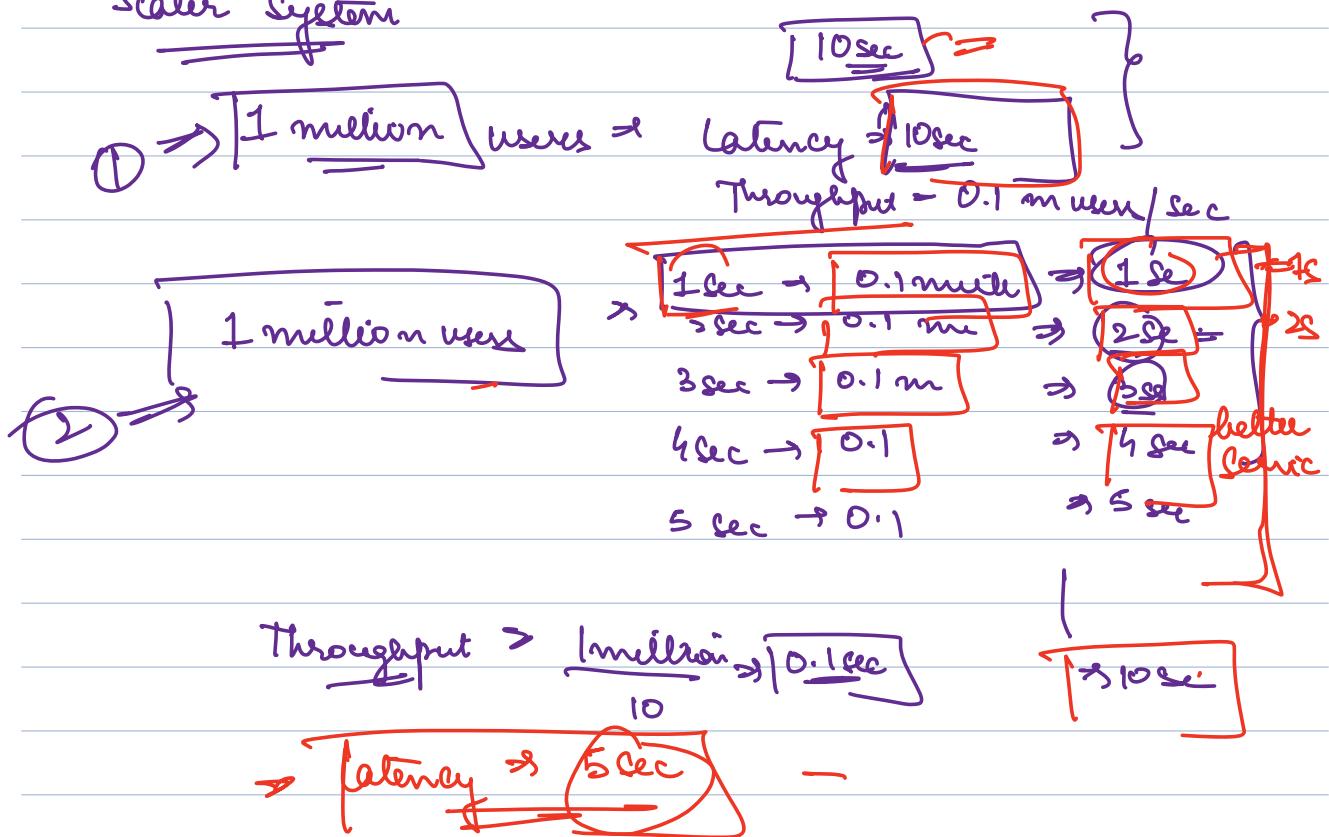
$$\overline{T_{100\text{proc}} \Rightarrow \frac{20 \text{ sec}}{1 \frac{100}{20} \text{ sp/sec}}}$$

Latency \Rightarrow Average time it takes for 1 process
to complete

- \Rightarrow time for 1 process \Rightarrow Completion time - Arrival Time
- \Rightarrow measure of perf of an algo

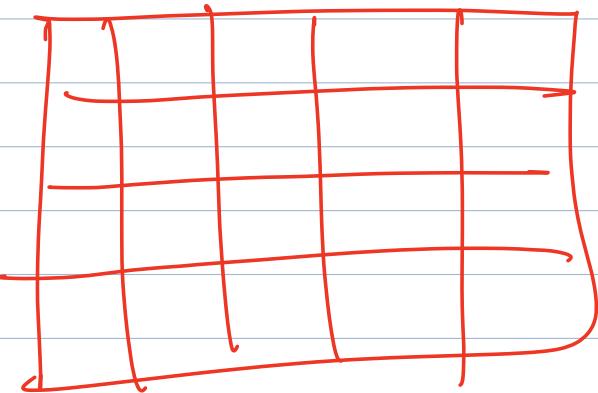


Scalar System



Eg

Restaurant



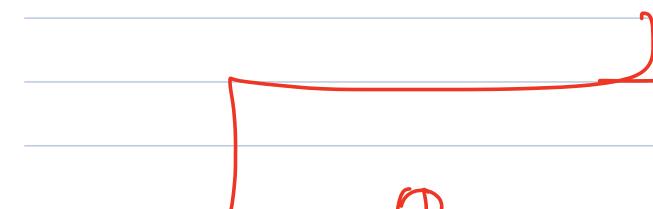
50 people

→ 50 people

→ all make order

→ To serve all → 50 min

⇒ Throughput $\frac{50}{50} \Rightarrow 1 \text{ min}$



I 0

at 50th min
→ it serves
order of everyone

II
0 5th → 10 min
15th = 5 order
20th → 10
50th → all

Lesser latency

5 orders
2 min 2 min 4 min 5 min 6 min

$$\frac{2+3+4+5+6}{5} \rightarrow$$

How to handle starvation

⇒ SJF ↪

⇒ SRTF =

⇒ Priority Scheduler ↪

result of priority

→ lower priority items are

prone to starvation.

| process | | arrival time | priority | comp |
|---------|---------|--------------|----------|------|
| i/d | process | | | |
| 1 | 1 | 1 | 4 | 4 |
| 2 | 3 | 3 | 2 | 2 |
| 3 | 4 | 4 | 1 | 3 |
| 4 | 5 | 5 | 5 | 1 |
| 5 | 5 | 5 | 6 | 2 |

highest priority

10

lowest priority
→ Starve

⇒ Aging ⇒ every time a CPU scheduler runs:

① it will select one process

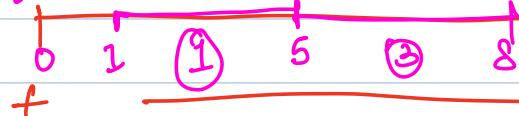
{ ② increase priority of every other process}

⇒ if multiple with same priority ⇒ [FCFS]

{ 1 } { 2 }

{ 3 }

{ 4, 5 }



+

+

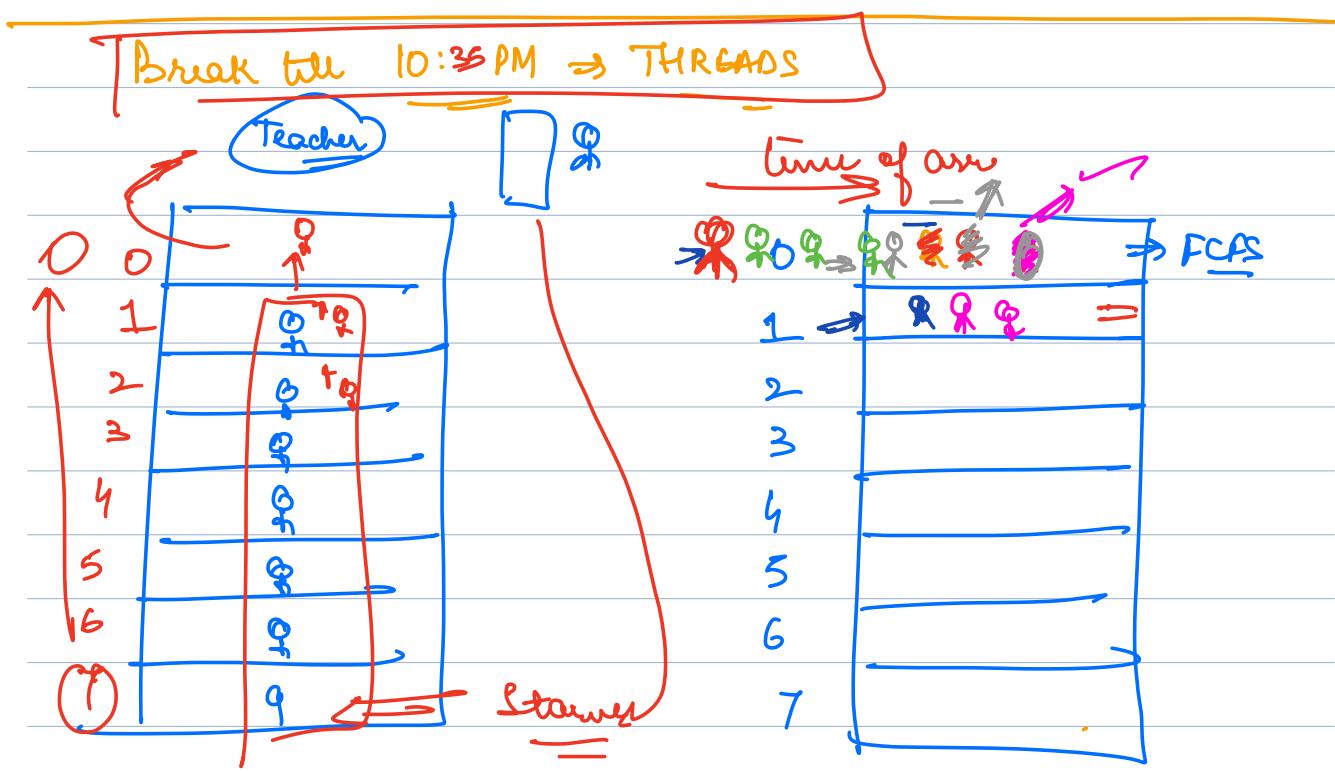
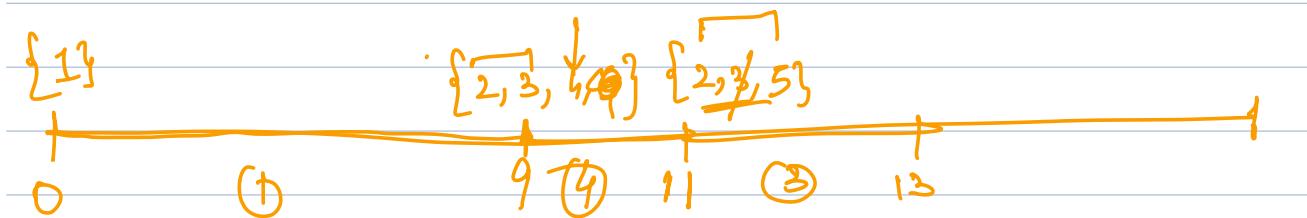
→

Upper bound

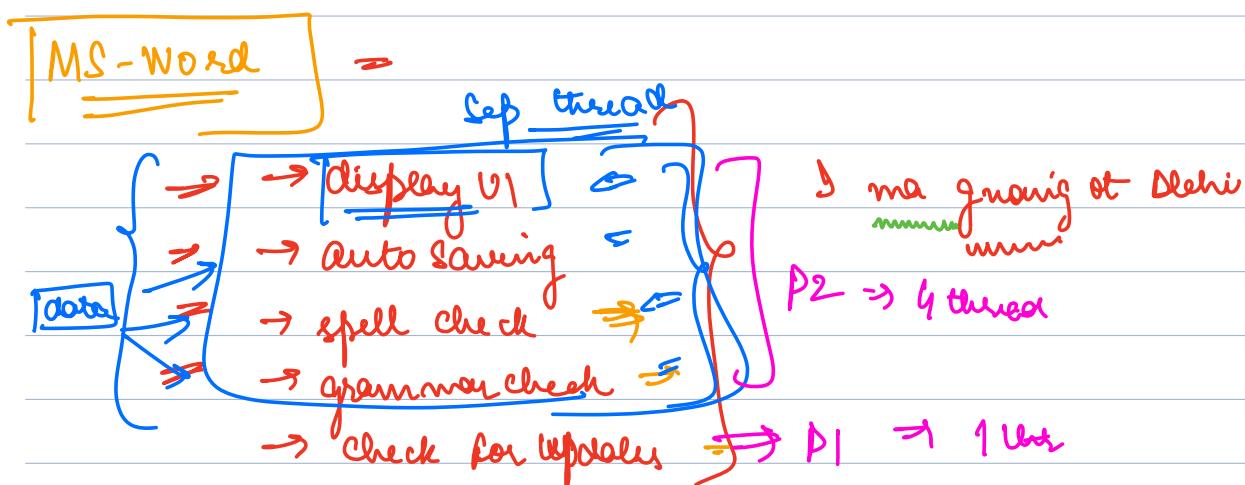
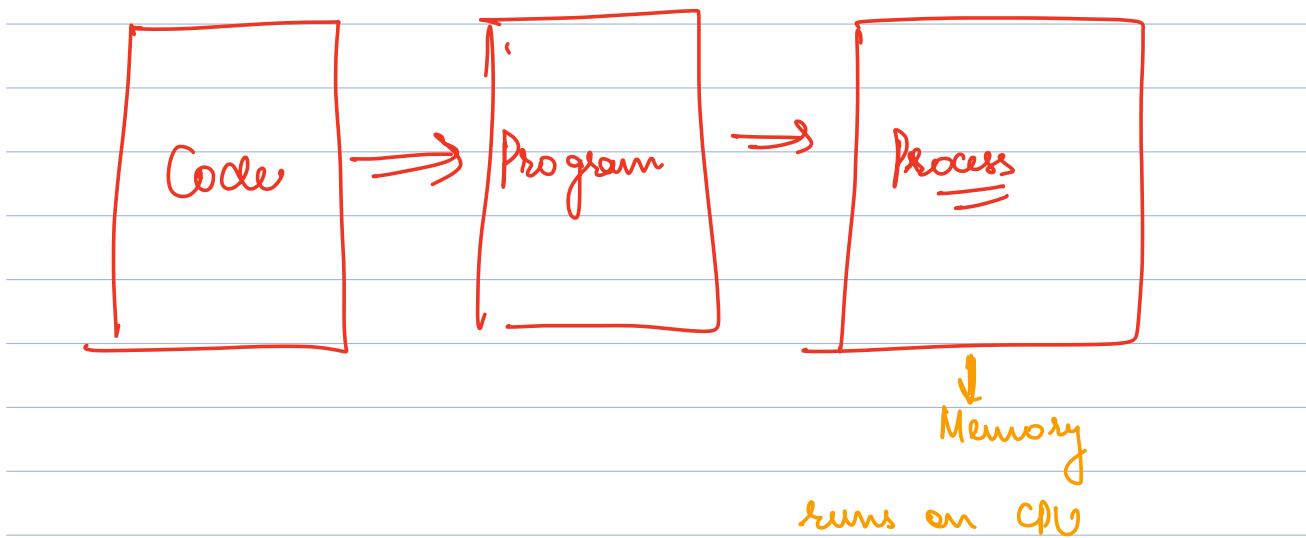
[priority | + | Order]

processes

| ID | arrival time | priority | Time to Comp. |
|----|--------------|-----------|---------------|
| 1 | 0 | 0 | 2 |
| 2 | 5 | 3 2 1 4 | 2 |
| 3 | 8 | 2 ① | 2 |
| 4 | 9 | ① | 2 |
| 5 | 11 | ② ③ ④ ⑤ 0 | 2 |



THREADS



Thread

- Basic / fundamental unit of CPU execⁿ
- CPU executes only threads
- whenever anything is running on CPU → it is a thread
- Process will atleast have 1 thread

main() {

```
print (Hello);
```

3

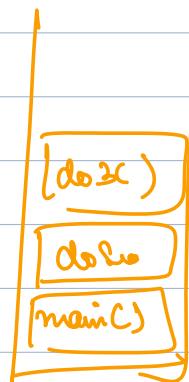
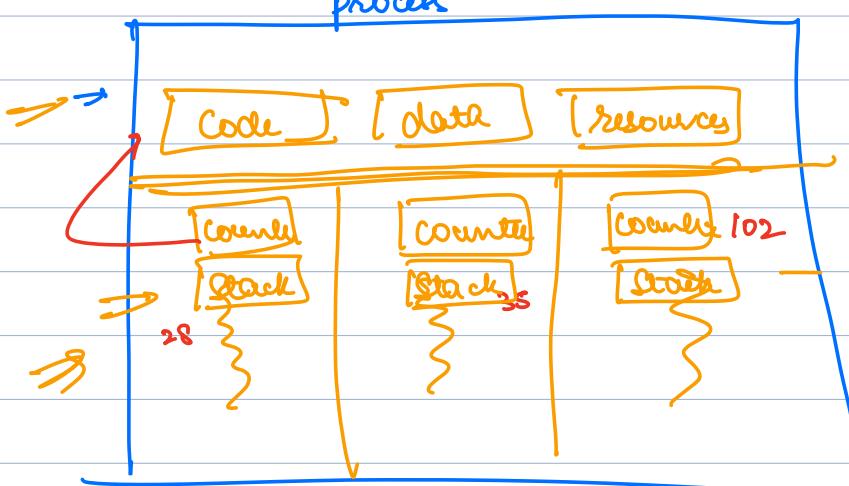
Process



next line
of code will
execute

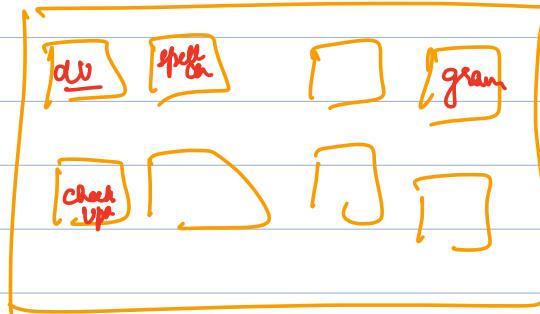
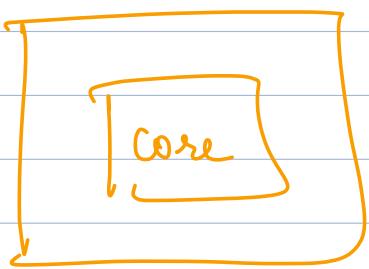
+ 1 thread

Because each process has its own data, data sharing b/w processes is difficult.



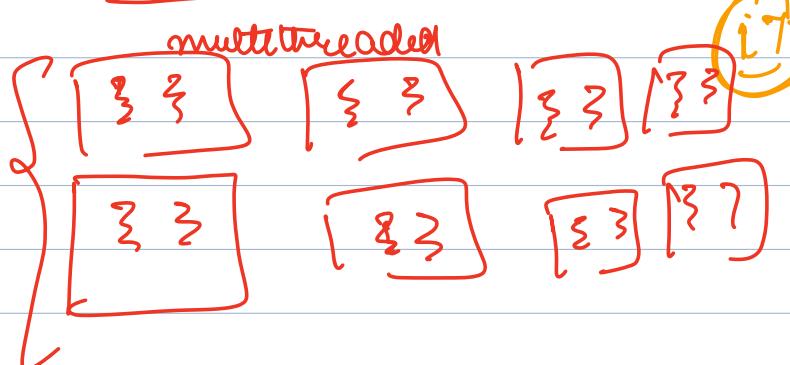
⇒ all threads share data of the parent process

⇒ every CPU core is executing a thread
⇒ thread scheduling



MS Word

O.S.
get spec class VI GC SC



8 core → 16 HyperThreading

CONCURRENCY vs PARALLELISM

CONCURRENCY

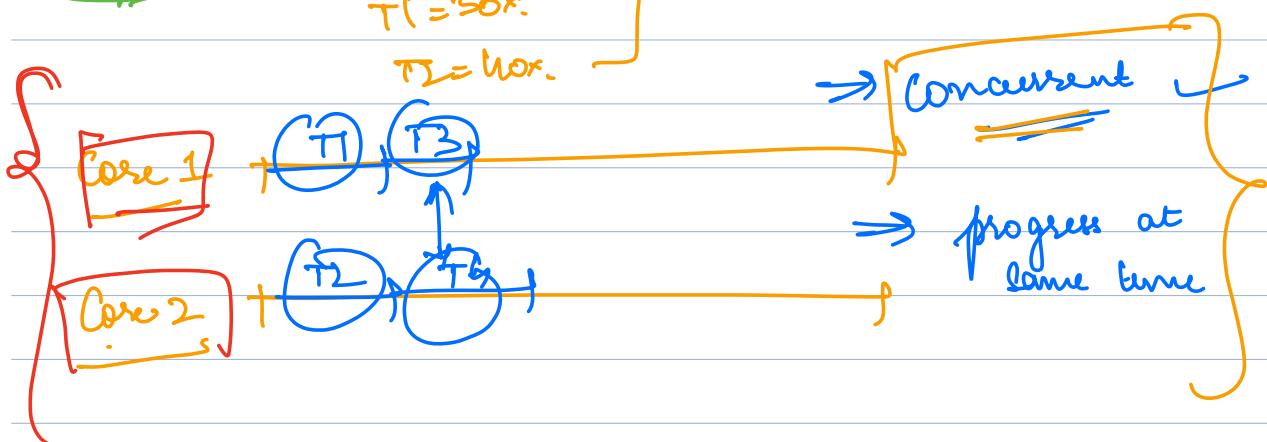
- Multiple threads at the same time in different stages of execution
- They may or may not be making progress at the same time



W/S serializable ⇒ tasks happen one after other



$T_3 = 30\%$
 $T_1 = 50\%$.
 $T_2 = 40\%$.



PARALLELISM

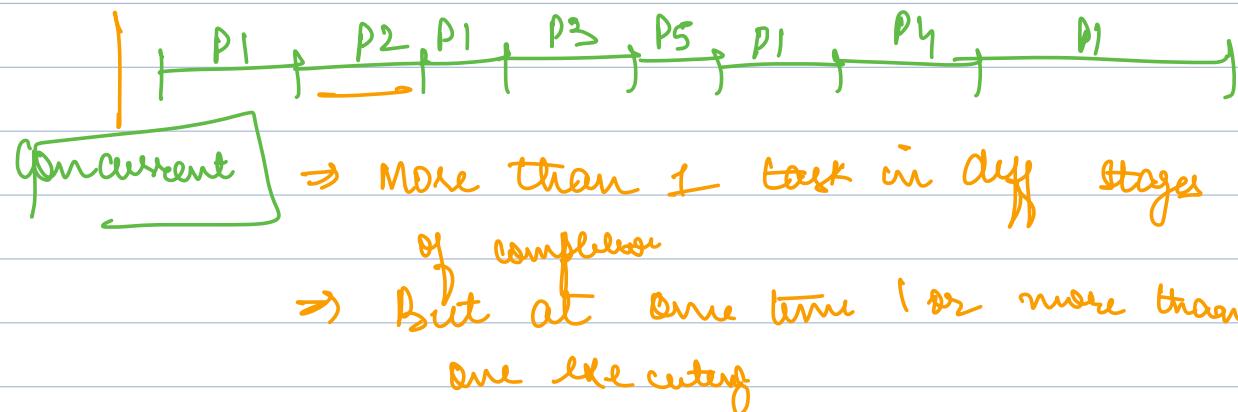
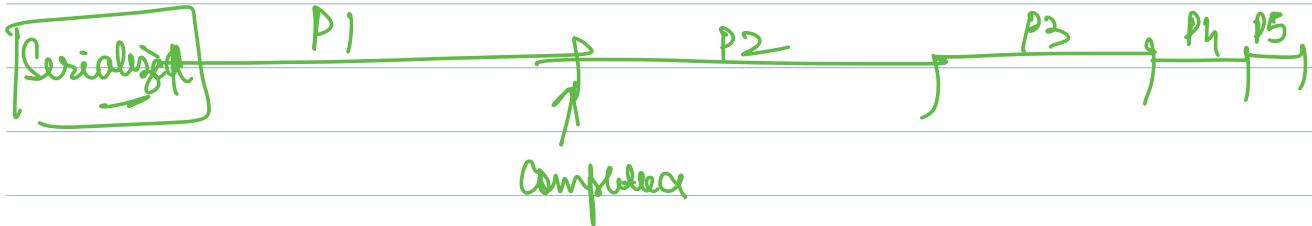
⇒ A system where more than 1 task can make progress at same time



{ Concurrent system may not always be parallel
Parallel System ⇒ always concurrent

Single Core ⇒ Can't be parallel
⇒ Can be concurrent

$[P_1, P_2, P_3, P_4, P_5]$



Parallel \Rightarrow more than 1 making prog

$$\begin{aligned} \text{CPU} &= 4 \text{ cores} \\ \text{GPU} &= 1000 \text{ cores} \end{aligned} \Rightarrow \begin{aligned} 2.8 \text{ GHz} &= \\ 400 \text{ MHz} &=] \text{Gauss} \end{aligned}$$

