

# Agenda

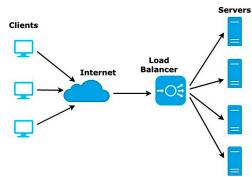
- 1. Sharding support@scaler.com
- 2. Load Balancer Routing +91 7351 769231
- 3. Consistent Hashing
- 4. Stateful vs Stateless servers

Delicious → bookmarking service

↓ simple → 1 computer.

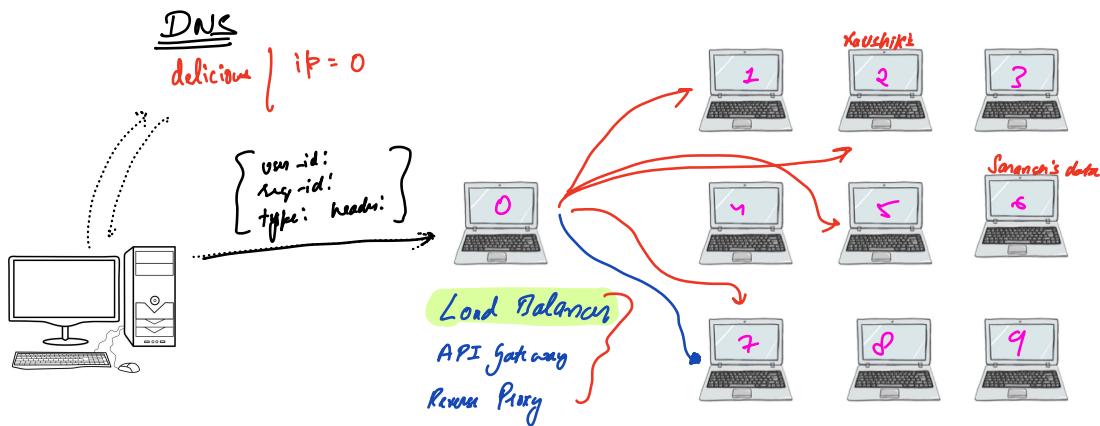
users ↑ ⇒ buy more expensive computers → Could not  
vertical scaling scale infinitely.

⇒ buy more of cheap servers  
horizontal scaling.



# Load Balancing

<https://docs.google.com/document/d/1DxQzLpu1XPemRWsewNWtKL6E4uwKHQhBp7GX6Sg7qI/edit>



❓ How does the Load Balancer track which servers are up & running?



Heart beat

Servers will push update to Zookeeper

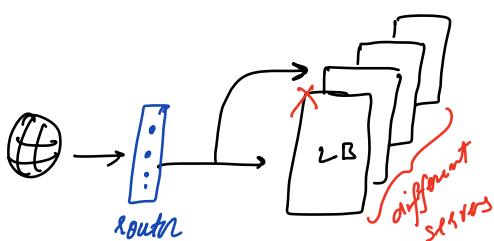


Zookeeper

Health check

Zookeeper will pull from servers

❓ What if the Load Balancer goes down?



DNS can also act as a LB on top of multiple LBs.

❓ Which machine should we send the request to?

Routing Algorithm

<https://docs.nginx.com/nginx/admin-guide/load-balancer/http-load-balancer/>

Goals

- ① Even Load Distribution
- ② Requests should not get dropped
- ③ easily add & remove servers (without large overhead)



## How to store the data?

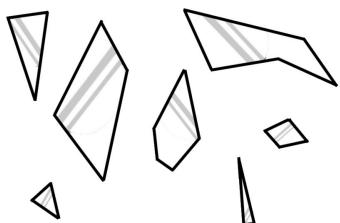
? Can we store all data on 1 machine?

No!

data is too large  
doesn't fit on single machine

? Split randomly?

challenge → how do we refine it??



## Data Sharding \*

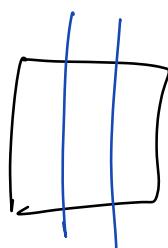
Split data across multiple servers horizontally according to some logic

Partitioning

split vertically  
↳ normalization

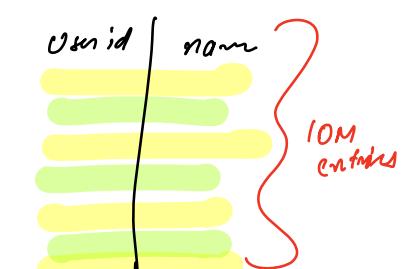
user-details

user-id	name	city	House no	Pincode
extract into separate table				



Sharding

→ data is too large  
→ improve both read & write performance at the same time



user-details-1

user-id	name
1	Alice

server 1

user-details-2

user-id	name
2	Bob

server 2



## Simple mod (Round-Robin)

$v_0 \rightarrow S_0$	$v_5 \rightarrow S_0$
$v_1 \rightarrow S_1$	$v_6 \rightarrow S_1$
$v_2 \rightarrow S_2$	$v_7 \rightarrow S_2$
$v_3 \rightarrow S_3$	$v_8 \rightarrow S_3$
$v_4 \rightarrow S_4$	$v_9 \rightarrow S_4$

Shared by user-id

↳ how to choose → later

# servers = 5       $S_0 S_1 S_2 S_3 S_4$

# users = 10

logic  $\Rightarrow [ \underbrace{\text{server-id}}_{\text{calculated}}, \underbrace{\text{user-id}}_{\text{req}}, \underbrace{\text{known}}_{\text{known}} ]$

$$\text{server-id} = \text{user-id \% 5}$$

Pros

- ① v.v. simple
- ② No memory req.
- ③ even distribution

LB

server-count = 5

```
fn route-request (r) {
    user-id = r.user-id
    server-id = user-id % server-count
    r.redirect (server-id)
}
```

Cons

old = red  $\rightarrow 5$  servers  
new = blue  $\rightarrow 6$  servers (+1)

$v_0 \rightarrow S_0$	$S_0$
$v_1 \rightarrow S_1$	$S_1$
$v_2 \rightarrow S_2$	$S_2$
$v_3 \rightarrow S_3$	$S_3$
$v_4 \rightarrow S_4$	$S_4$
$v_5 \rightarrow S_0$	$S_5$
$v_6 \rightarrow S_1$	$S_0$
$v_7 \rightarrow S_2$	$S_1$
$v_8 \rightarrow S_3$	$S_2$
$v_9 \rightarrow S_4$	$S_3$

Server goes down

↳ data??

parent don't loss  $\rightarrow$  replication

↳ diff class

assume: central backup of all data at all times

With  $\rightarrow$  some load of old servers gets moved to new servers

what are we get moved?

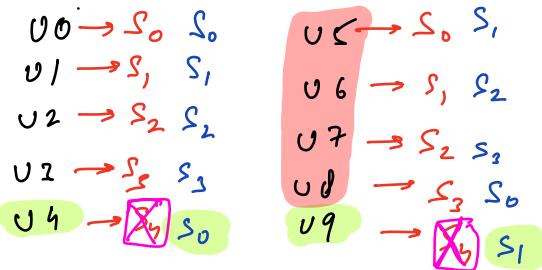
new user  
showing load

$v_5 \rightarrow 0 \rightarrow 5$
$v_6 \rightarrow 1 \rightarrow 0$
$v_7 \rightarrow 2 \rightarrow 1$

$v_8 \rightarrow 3 \rightarrow 2$   
 $v_9 \rightarrow 4 \rightarrow 3$

need less shuffling

old  $\rightarrow$  5 servers

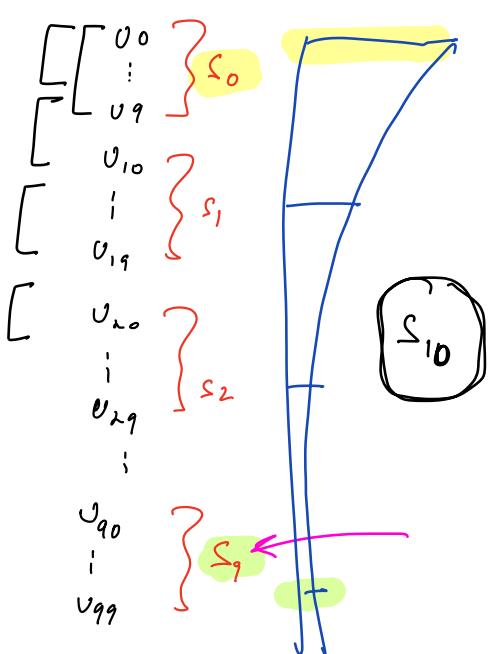
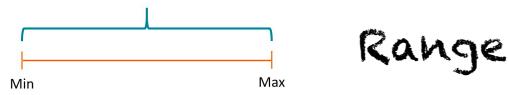


new  $\rightarrow$  4 servers

S4 crashed

aim: data earlier stored at crashed server ( $S_4$ )  
gets mapped to other servers.

some users get mapped for no reason  
and less data shuffling.



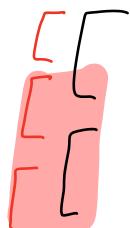
Pros ① Simple

Cons

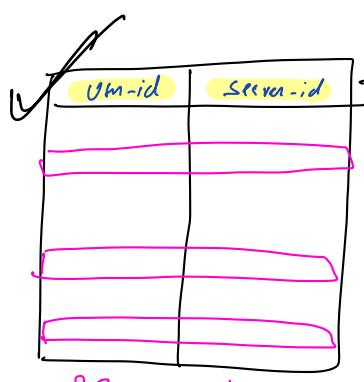
① Equal distribution?

to add new users, you need to add new servers

② Add / remove X



## Mapping Table



→ LR

Pros

- ① equal distribution
- ② crosses 3
- ③ added-

Cons

large memory req.

12 B for entry,

2 B over

⇒ needed memory  
= 12 bytes + 2 B ⇒

24 GR  
RAM

→ fast

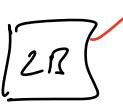
HDD → slow



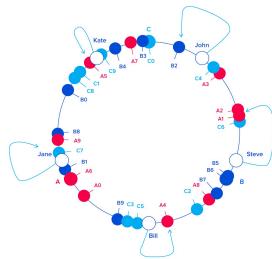
Sync!



Sync!



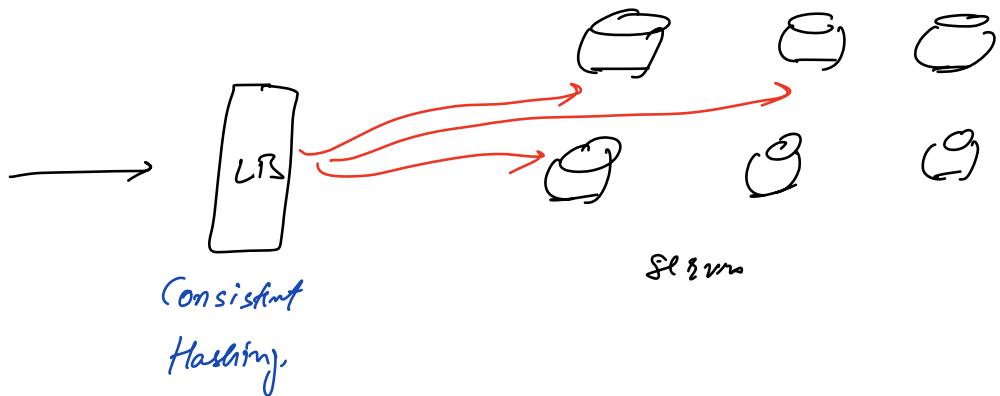
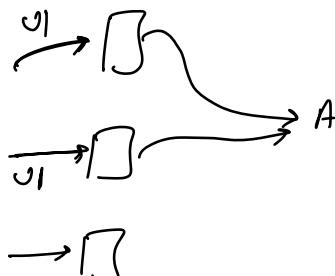
Sync! → v.v.v.v  
v. difficult



# Consistent Hashing

- Pros

  - ① Simple & fast ✓
  - ② Does needs v. less memory (RAM) ✓
  - ③ Can be distributed (no sync needed) ✓
  - ④ fairly add & remove servers ✓
  - ⑤ equal load distribution



**Hash** → function that takes in 1 value & returns  
another value  
 $\text{input} \rightarrow \text{anything}$   
 $\text{output} \rightarrow \text{well defined range}$

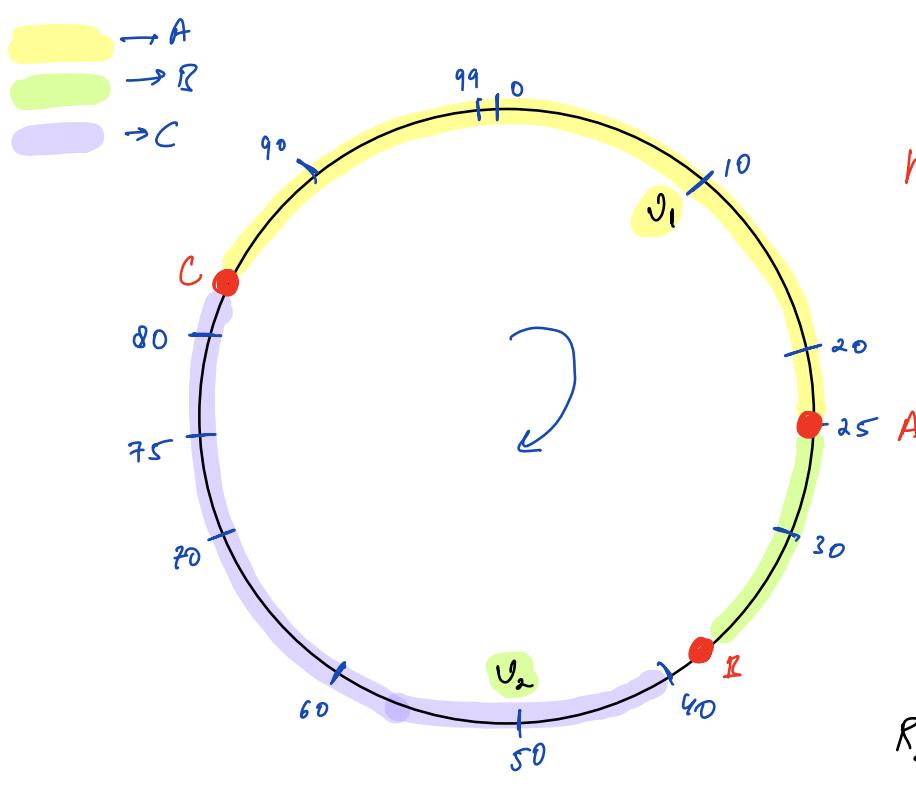
$$h_1(\text{string}) \rightarrow \left( \sum_c \text{ascii}(c) \right) \% 12$$

$$h_1 \left( \begin{smallmatrix} a & b & c \\ 1 & 2 & 3 \end{smallmatrix} \right) = \left( 97 + 98 + 99 \right) \% / 13$$

Possible output  $\Rightarrow$  0 - 12

$h_1, (-)$  }      64 bit intgn       $(0 \rightarrow 10^{18})$   
⋮  
 $h_{10}, (-)$  }      range

$$P_{\text{collision}} = \frac{1}{10^{18}}$$



Servers  $\Rightarrow A \quad B \quad C$   
 $\text{hash}(\text{IP}_A) = 25$   
 $\text{hash}(\text{IP}_B) = 37$   
 $\text{hash}(\text{IP}_C) = 84$

### Requests

$R_1 \rightarrow v_1 \rightarrow A$

$$\text{hash}(\text{user-id})$$

$$h(v_1) \Rightarrow 10$$

$R_2 \rightarrow v_1 \rightarrow A$

$$h(v_1) = 10$$

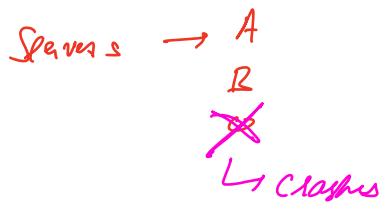
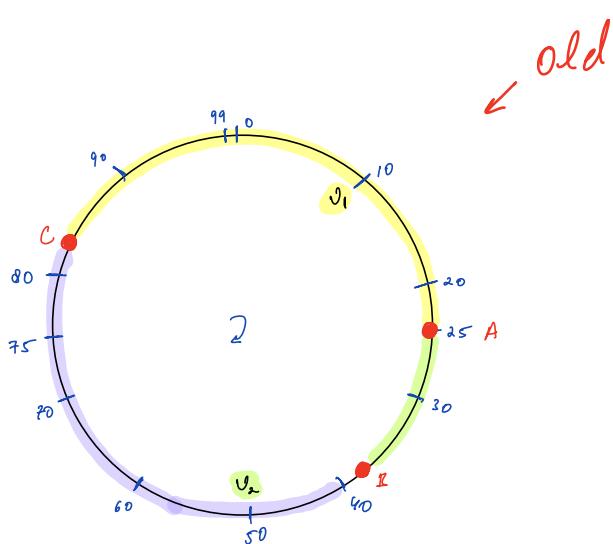
$R_3 \rightarrow v_2 \rightarrow C$

$$h(v_2) = 50$$

- ① visualize hash output as a ring
- ② hash each server & placing on ring
- ③ given a req, find user-id from req. hash it.  
find nearest server going clockwise.

④ User  $i$  always goes to same server (if server is up)

## Server Crash



result →

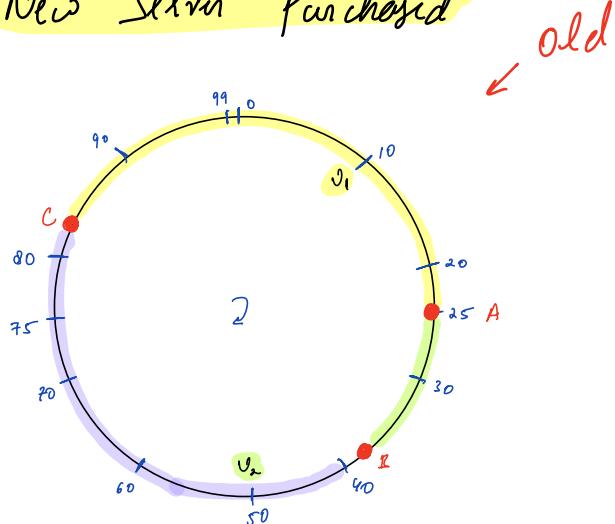
All users earlier mapped to  
Server C now get  
re-mapped to server A.

data move

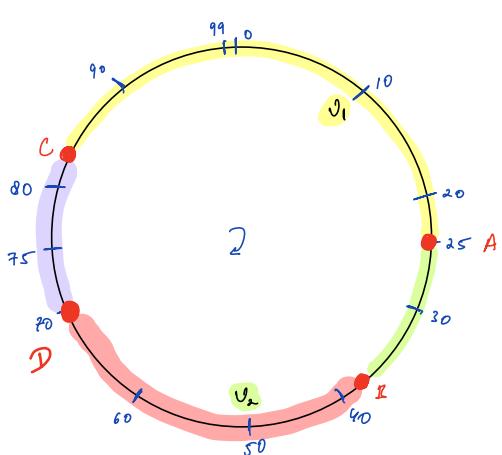
backup  $\xrightarrow{C} A$

Unnecessary data movement?  $\times$

New Server Purchased



Servers  $\rightarrow$  A  
B  
C



add Server D  
 $hash(ip_D) \Rightarrow 70$

Some users earlier mapped to C  
are now mapped to D

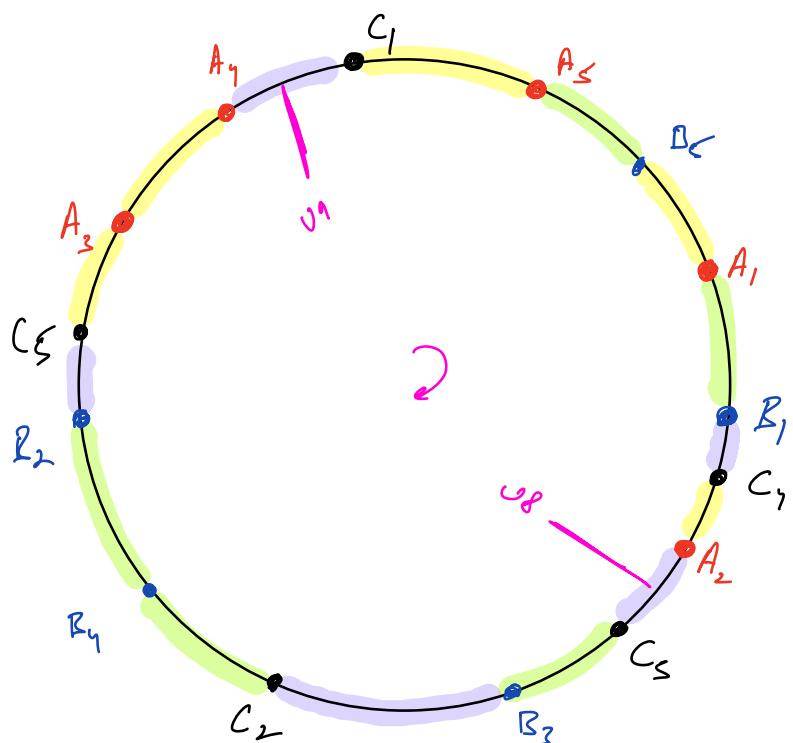
no unnecessary data shuffling.

$hash(1) \rightarrow$    
 $hash(4) \rightarrow$

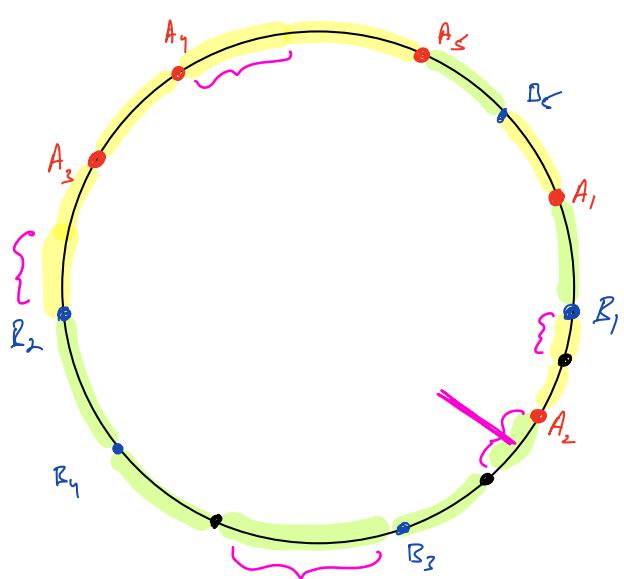
Magic

- ① Use  $\lambda$ -hash functions to assign multiple locations to each server.

- ② User still gets hashed only 1 time.

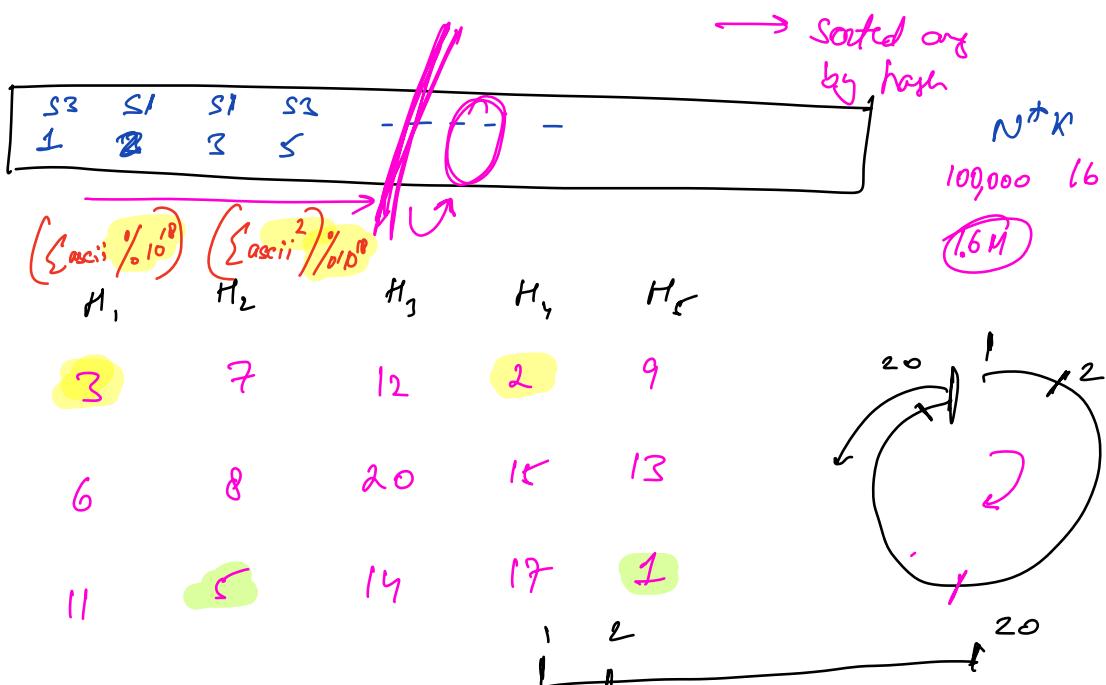
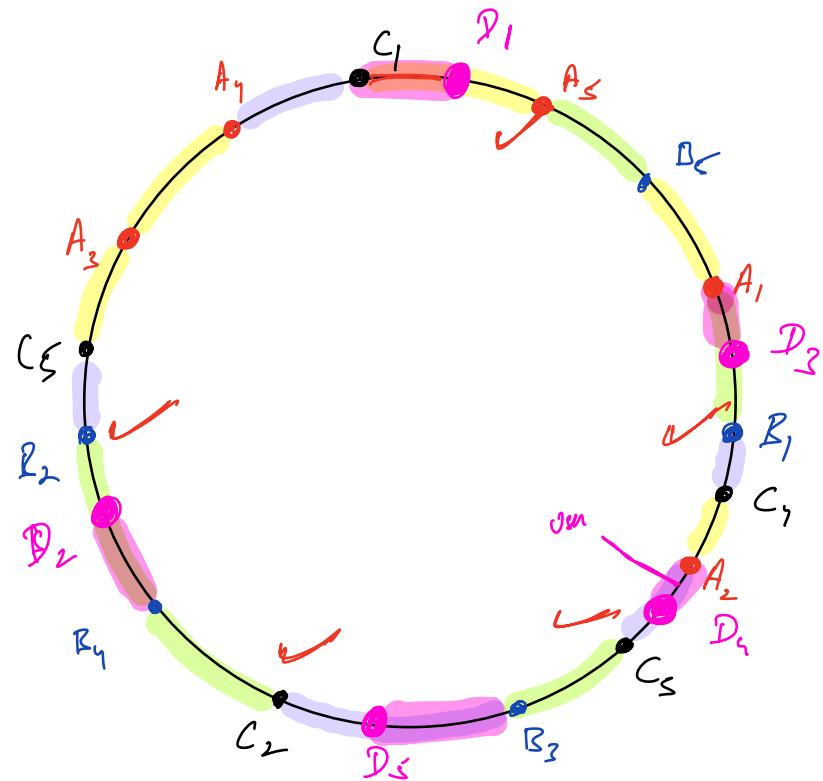


Server  
C Crashes



Server C's load  
was equally distributed  
across remaining  
servers.

## Purchase Segments D



$$H_0(j_i) = \underline{15}$$

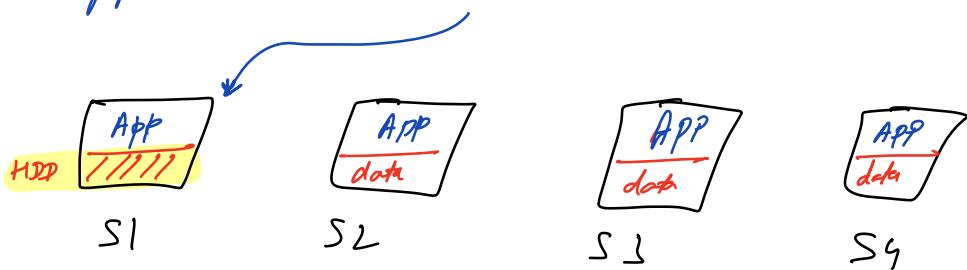
$\sum_{i=1}^{N^{\frac{1}{2}K}} (\text{ASCII}_i)^2 / 10^{16}$



## Stateless Servers

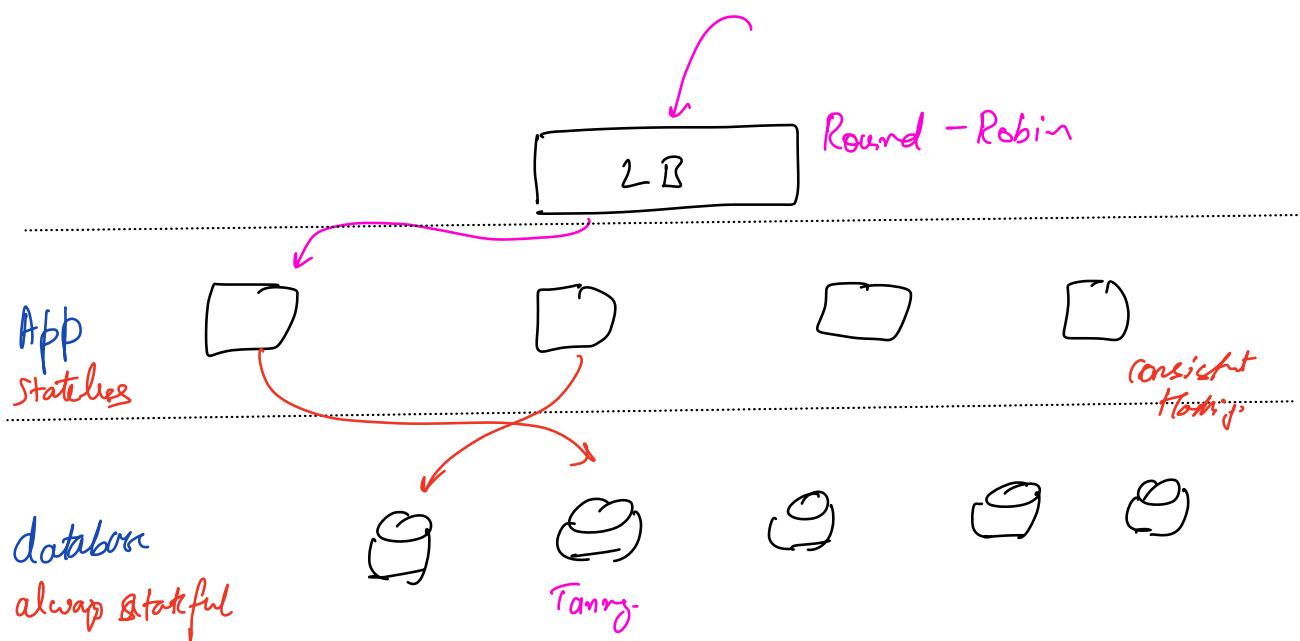
any data that needs to be stored.

Both app & data live on each server.



## Stateful servers

State → Consistent Hashing.





Fannay (read)

LIP

randomly assign



consistent hashing



App  
Layer



Rahul



Sonam



Dashan



Tam



Vivek

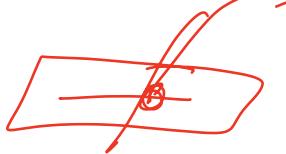
Database  
Layer

Select \* from order  
where order\_id = req. order id

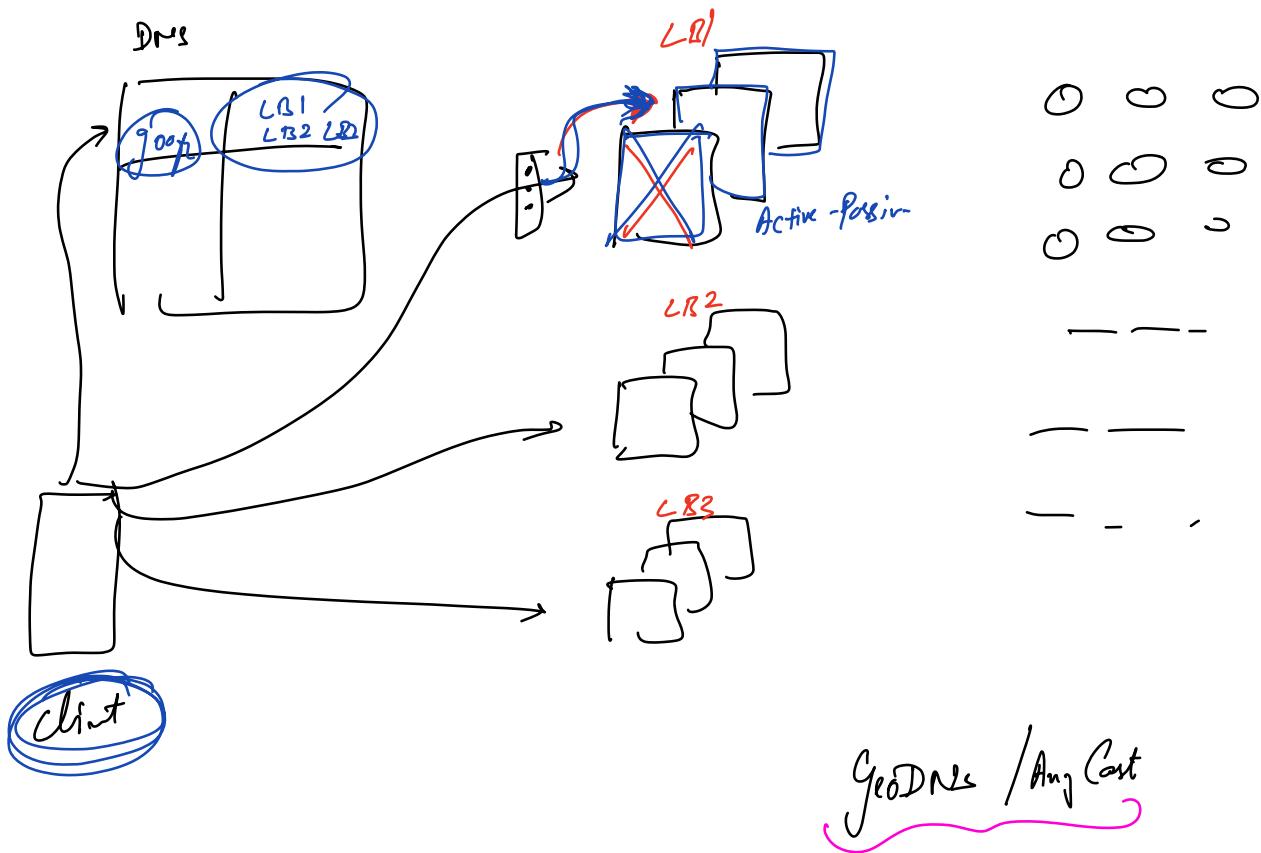
20

hash(20)

= v

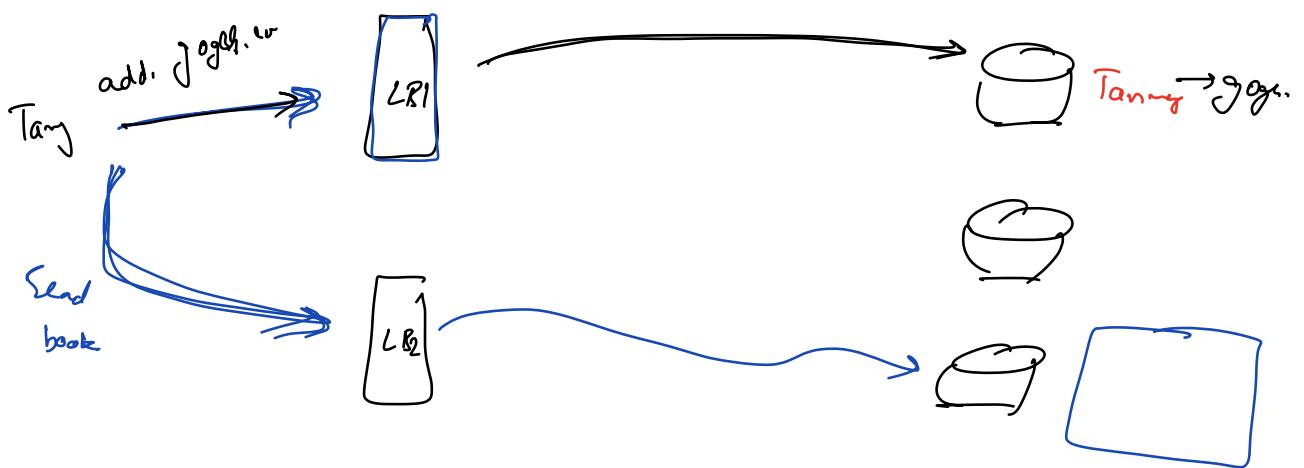


DNS as a LR → Balancing Load  
not for fault tolerance



Prod - market fit

CTO  
bias of action



$v_1 \rightarrow LB_1 \rightarrow S_1$

$v_1 \rightarrow LB_2 \rightarrow S_2$

