

# Agenda

1. Appserver Layer
2. Caching
3. Invalidation & Eviction
4. Local Caching

support@scaler.com

+91-7351769231

Scale is large

↳ Read Load → 90% app, **Read-heavy.**

↳ Write Load  $\# \text{reads} \geq 10^6 \text{ Harrys}$

→ Caching helps us handle a large Read-Load

# Writes  $\leq$  # Reads (99.99% cases)

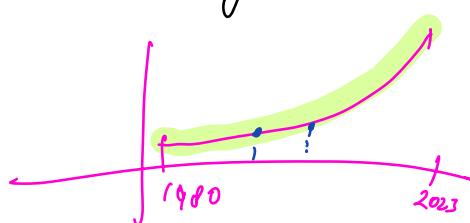
# Writes  $>$  # Reads

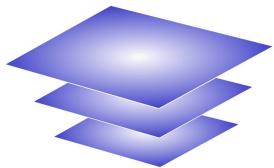
data goes in, but doesn't come out!!

↳ archival

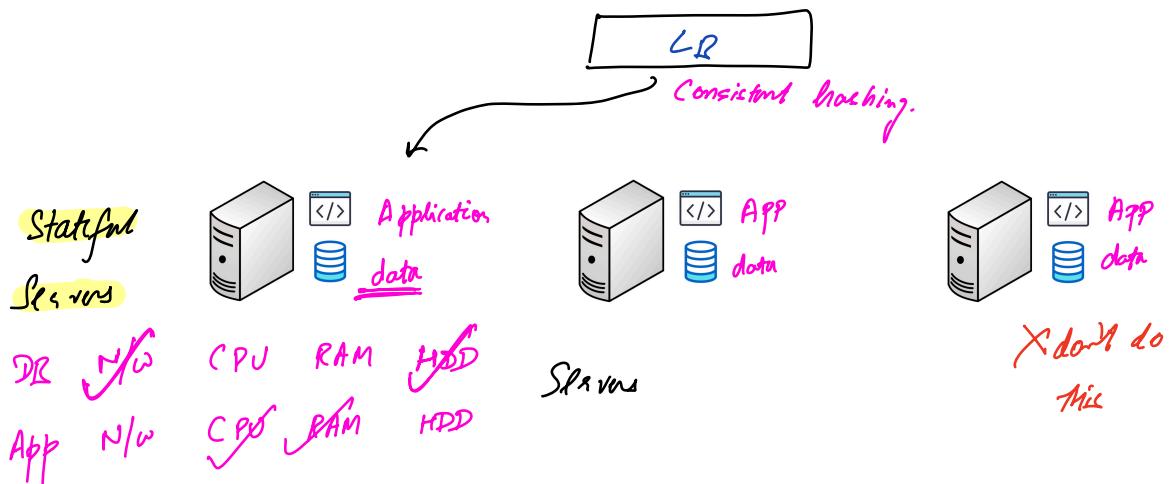
↳ logs

↳ analytics



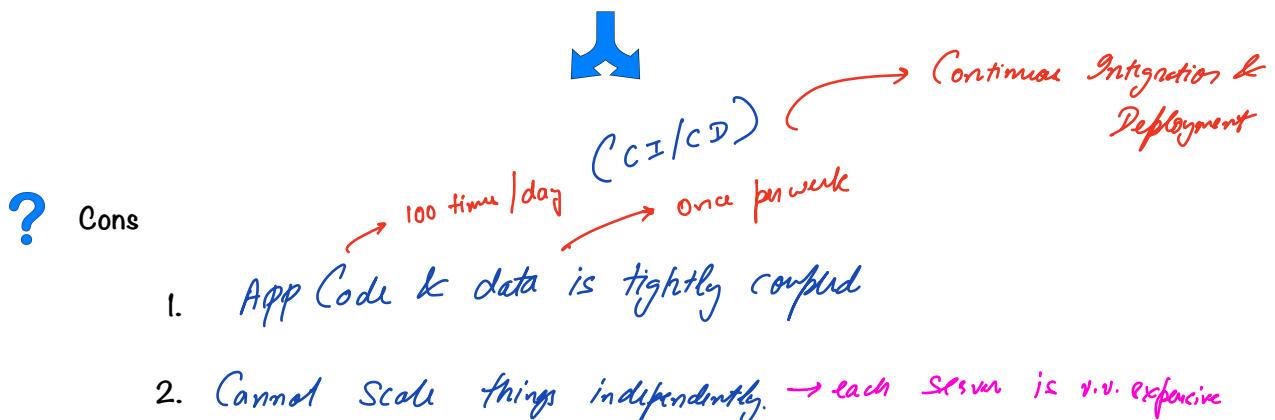


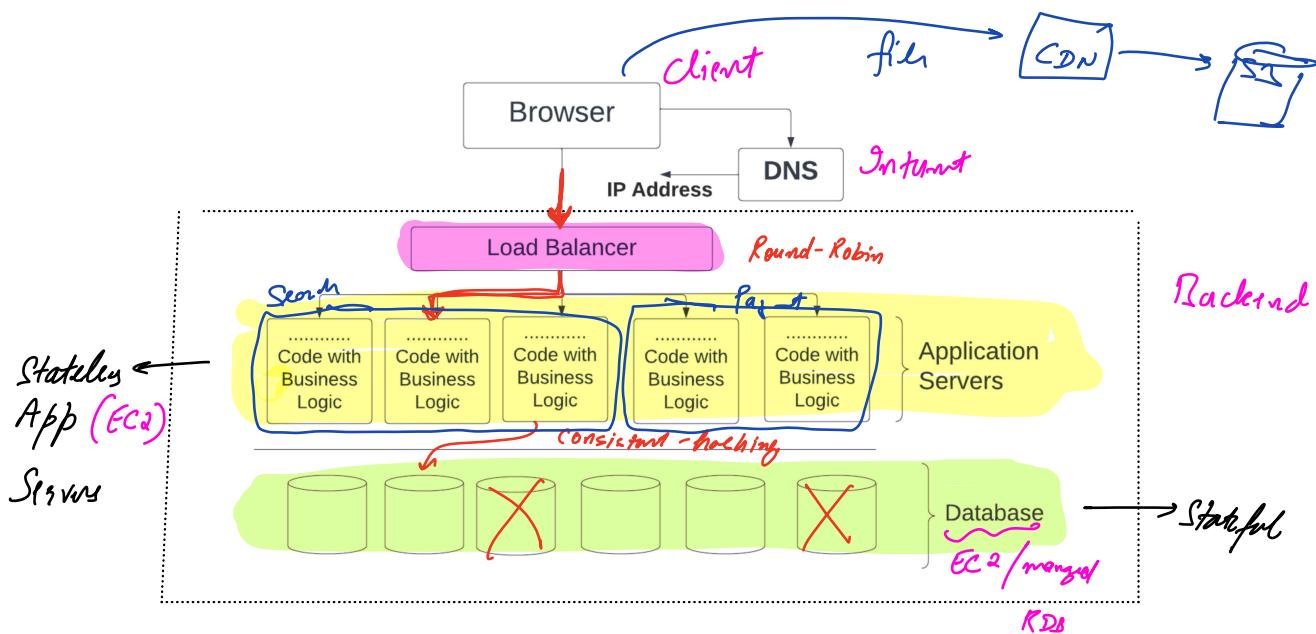
## Appserver Layer



? Is this good design?

1. Data access is fast ∵ data is local.
2. Simple to code & deploy.
3. Transaction/monitoring/security /testing → easier.





❓ Which request to which server?

doesn't match → Round Robin

∴ app servers are stateless



Who does the Consistent Hashing? Load Balancer / App Server

DR client library / DR orchestrator



← Orchestrator

ZooKeeper → keeps track  
which servers  
configuration

How do the Appservers know how many & which DB servers are available?

DR orchestrator will keep track of these



# Making Tea

?

Ingredients

Water

Tea leaves

Milk

sugar

ginger

Lemon grass

gloves

Source → industry  
for you → store / shops

Cardamom



Super Market



Kitchen / Refrigerator  
(Temporary storage)

→ Source

→ derived storage

→ all of the stuff is stored here

→ typically much smaller

→ up-to date / fresh stuff

→ can have stale items  
get rid of them

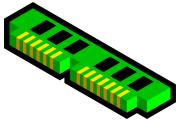
→ super slow / high latency

→ ultra-fast

→ always needs to be persisted on HDD

→ can be in HDD / RAM

Redis / Memcache / ECache



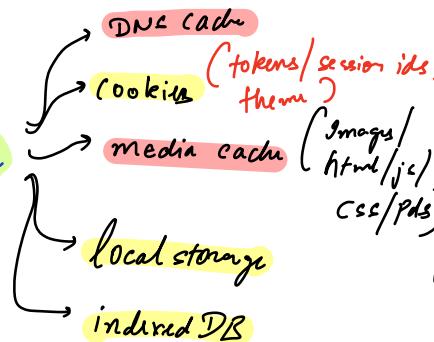
# Caching → always layered

client



## Browser Cache

files → CDN  
API req.  
↳ API



Computer Memory:  
HDD → level 3

RAM → level 2

CPU → Processor

500x ↓

L2 cache

L3 cache

L1 cache

Registers

flip-flops



0.0 ns



## Content Delivery Networks (CDN)

highly distributed

Cloudflare

Akamai

Amazon Cloudfront

Fastly

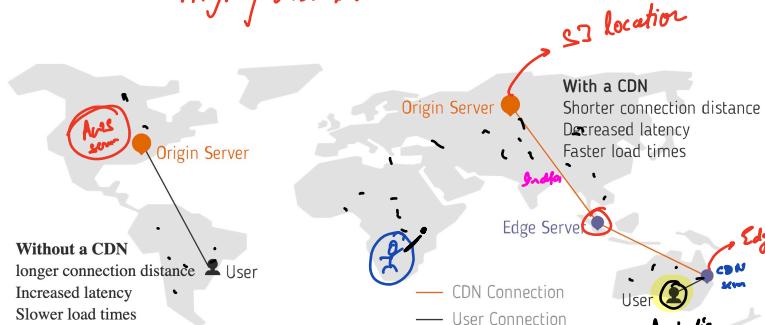
70%

all index

is served

via step 2

Edge  
↓  
close to user



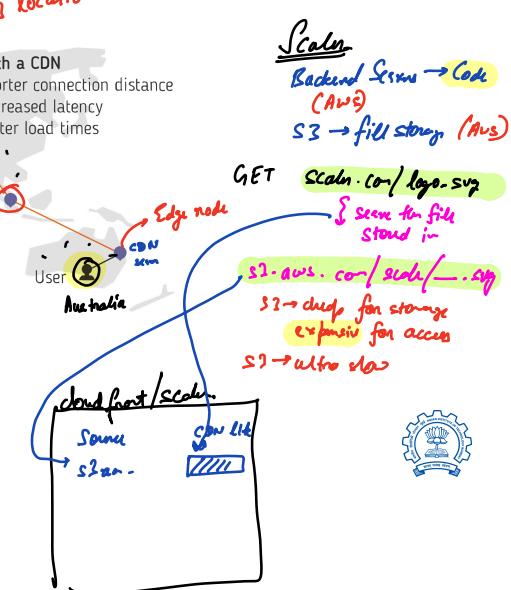
<https://www.cloudflare.com/learning/cdn/what-is-a-cdn/>

? How to get the IP of the nearest CDN server?

① GeoDNS → low penetration

② AnyCast → hard/trick for CDNs

↳ try to act as their own Geo DNS

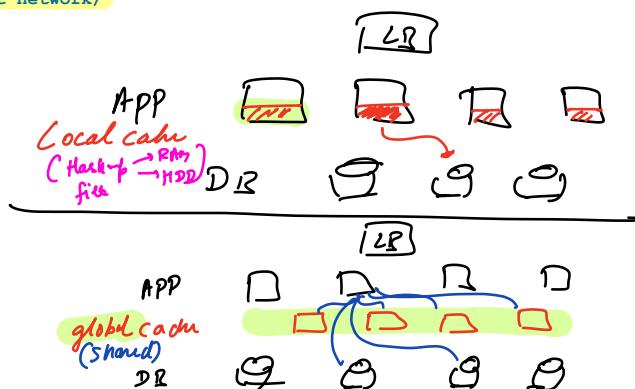


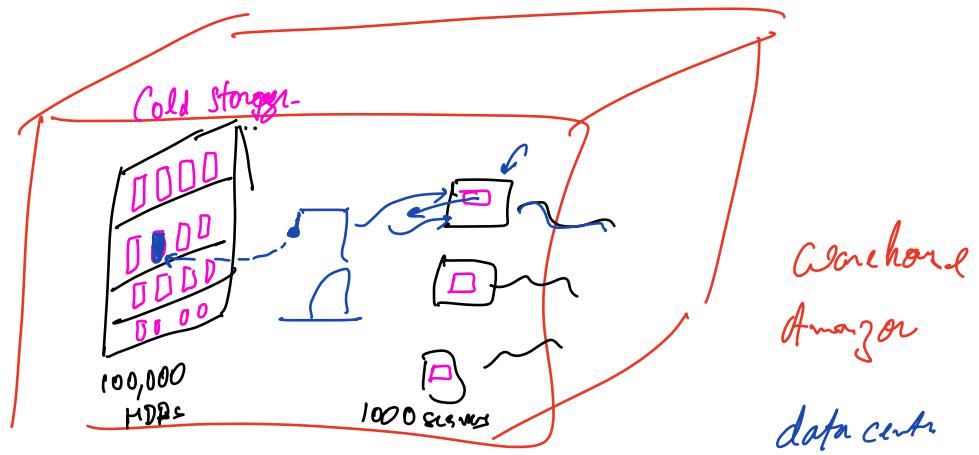
<https://www.cloudflare.com/en-gb/learning/cdn/glossary/anycast-network/>  
<https://constellix.com/news/anycast-vs-geodns>

→ backend ↓

Local Cache

Global Cache





10:17 → 10:30



# Challenges with Caching

Stale Data



Cache is not the  
source of truth → DB

data in DB is updated  
data in cache is not updated

Tiny



Cache can become  
full

to add new data  
↳ remove some other  
data (Eviction)



## Cache Invalidation

⌚ TTL (time-to-live)

→ Cache can have stale entries

→ Some of the cache might be unused

① Eventual Consistency

fast reads (10,000/s)

③ fast writes (100/s)

⌚ Eventual Consistency.

90%

## Write Through

① Immediate Consistency → No stale data whatsoever!!!

② fast reads (10K/s)

③ slow writes (10/s)

## Write Back

① No Consistency → data loss!!!

② fast reads (10K/s)

③ ultra-fast writes (10K/s)

## Write Around (v. similar to TTL)

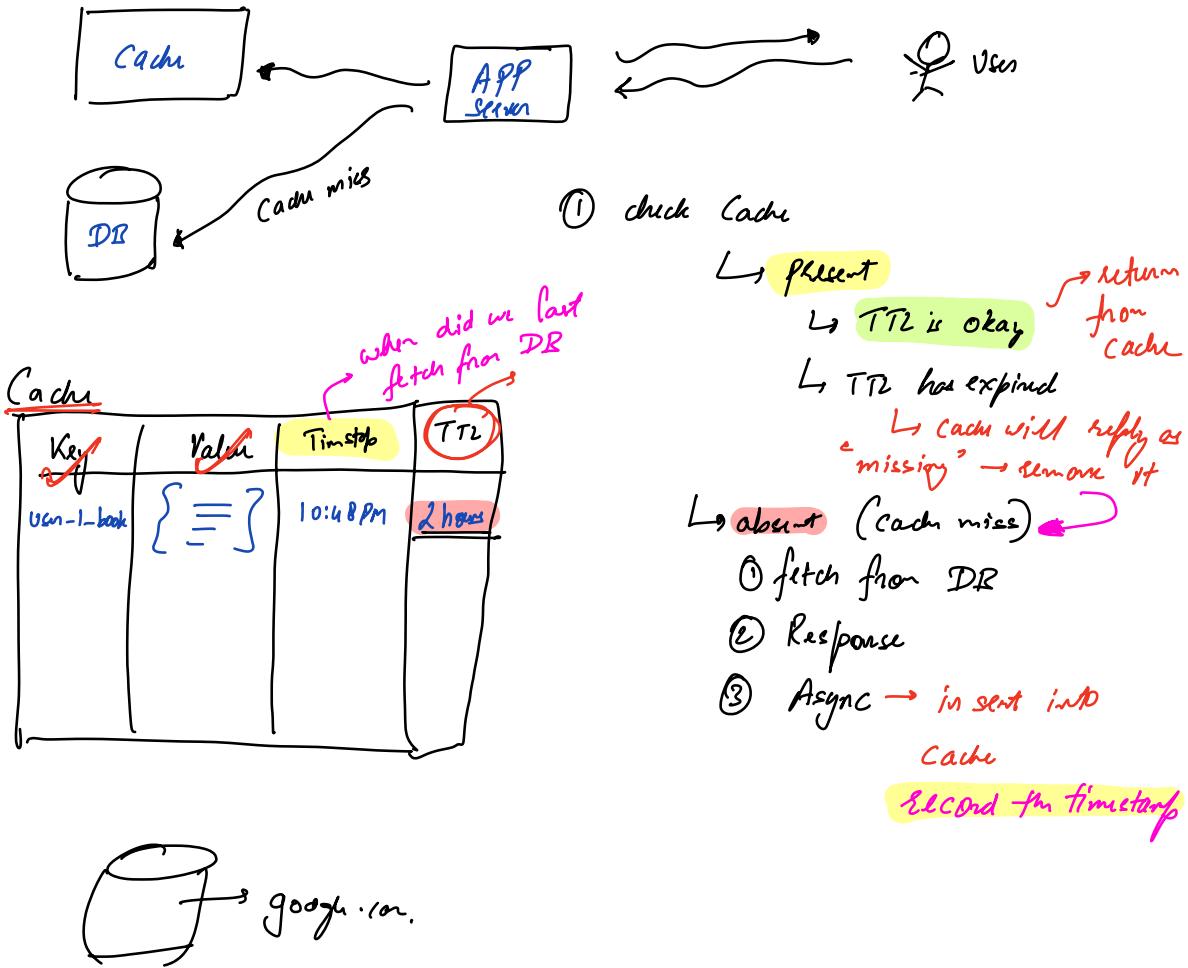
① Eventual Consistency

fast reads (10K/s)

③ fast writes (100/s)

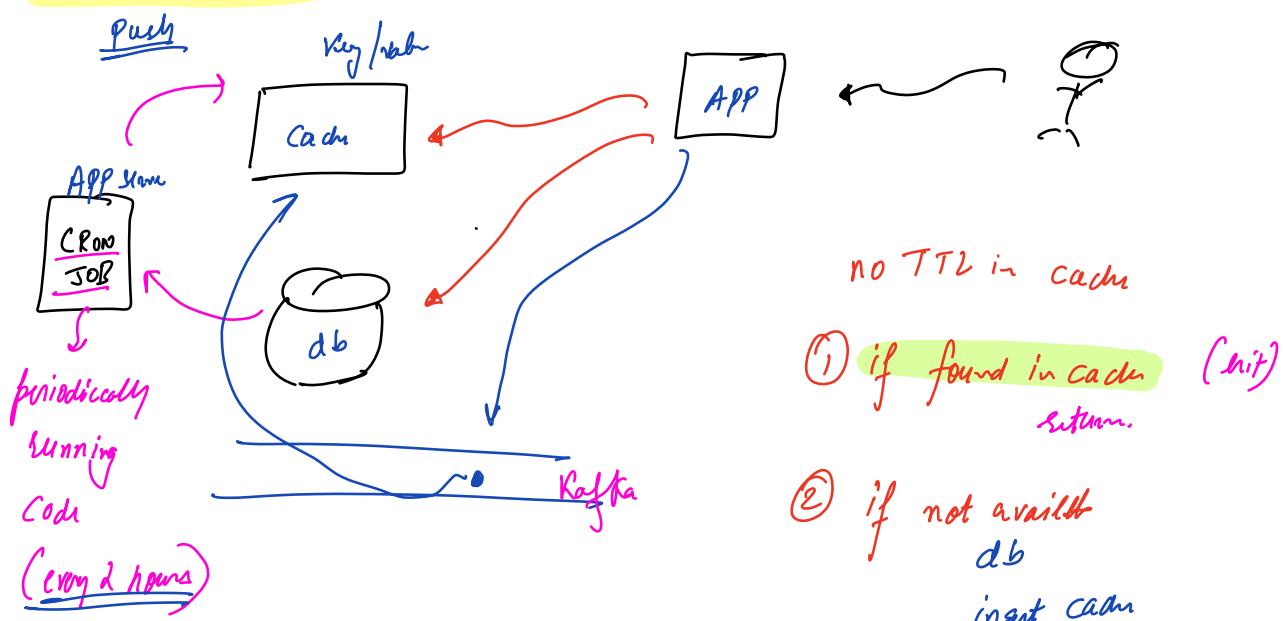
## TTL Pull

Time to Live  $\Rightarrow$  for how much time do we consider the data as sufficiently fresh (usable)



Write  $\rightarrow$  just to db

## Write-Around

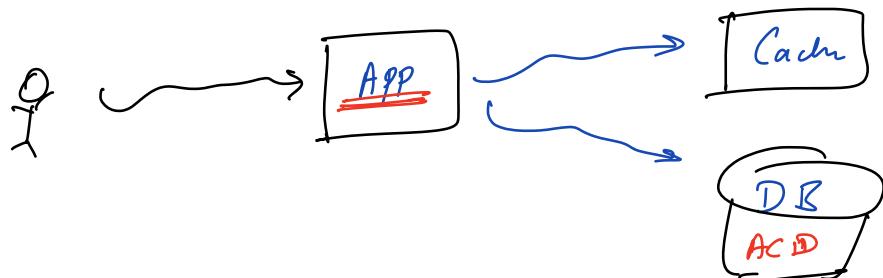
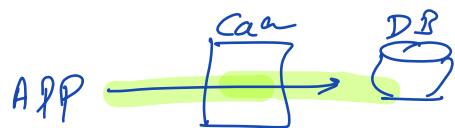


CRON

Select + from bookmarks  
where updated-at > (now() - 2 hours) } → update in cache.

Writely → just to db

## Write-Through



① check cache

↳ present → return

↳ absent

↳ read from DB

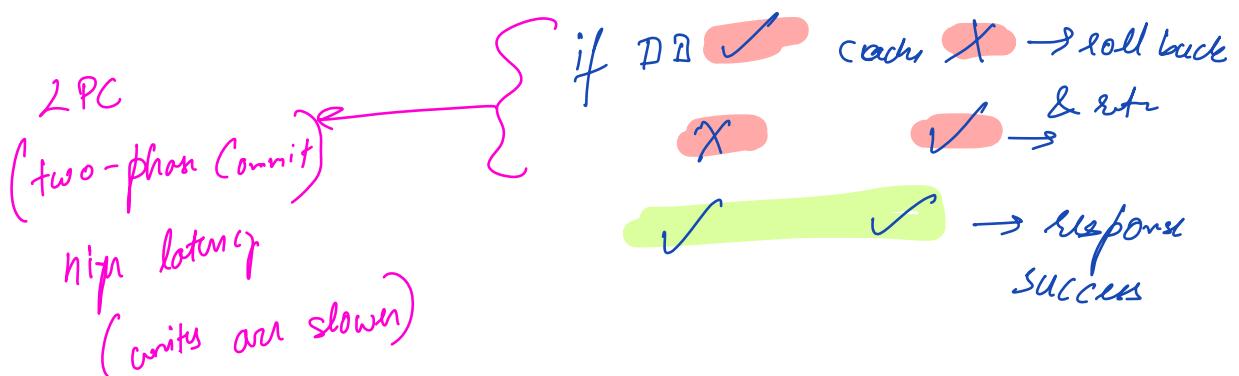
↳ insert into cache

?

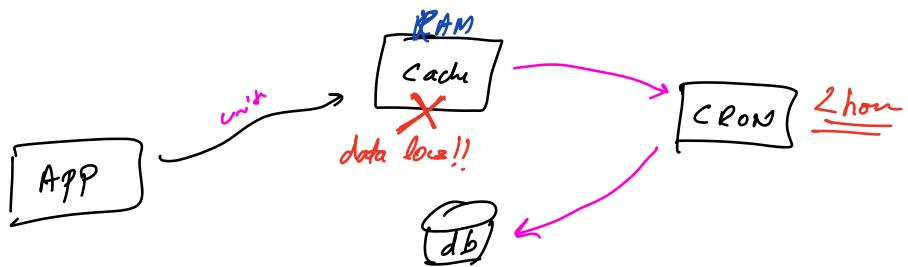
Read

## Writes

writes to both DB & Cache → transaction <sup>atomic</sup>



## Write-back



Write → simply write to cache

if cache crashes before sync → data loss

Youtube → view count      12/m / 100m / 10K  
Twitter →



## Cache Eviction

Similar to  
Page Replacement  
Policies

Strong Assumption: If a piece of data has been accessed

recently/frequently in past → more likely to be accessed in near future.  
↳ temporal locality

First in First Out  
Singer Queue

— not that good

Least Recently Used  
Priority: timestamp

89 %

Priority Queue (min-heap)  
LL + Hash Map → puffed

Last in First Out

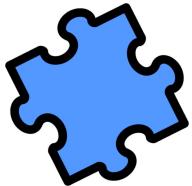
→ v. v. rare → special case → H/w

Least Frequently Used  
Priority: counter

9 %

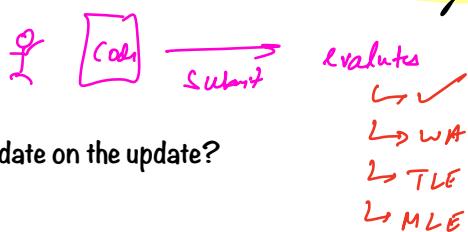
→ Similar to LRU  
implementation.

Social-shifty-website  
-that-moon-uuu.com/  
↳ 3 visitors in  
past 10 years.



## For next class

Code Judge



Local caching of test data on the app server. How to invalidate on the update?

300,000 people live

Many cache



How is facebook's newsfeed calculated, and how can you make it very fast?

2B users

→ updated freq.  
↳ read freq.

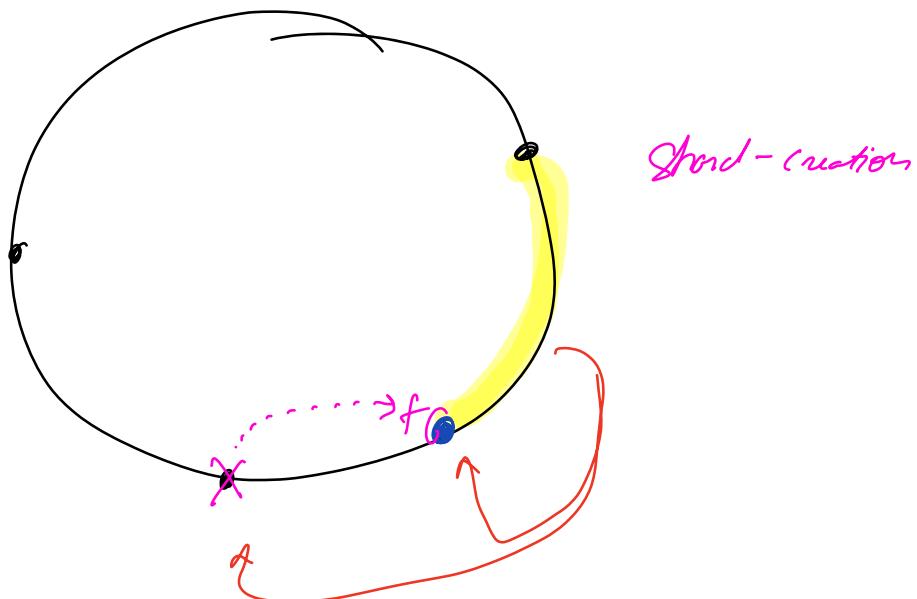


Unix → daemon process  
↳ background

ls → list directory  
mv → moving

Cron → periodic tasks

Cron job → periodic task



## Any Cast

